

Introduction to SQL



Morning lecture - July 19, 2017

Data Science Immersive, Galvanize Platte



General Objective



By the end of this lecture, you'll be able to connect to a database from the command line and use SQL to answer questions about the data.

Specific Objectives



- Discuss RDBMS and why we use them
- Write simple SQL queries on single table using SELECT, FROM, WHERE, GROUP BY, ORDER BY clauses as well as aggregation functions (COUNT, AVG, etc.)
- Understand primary keys, foreign keys, and table relationships
- Write complex queries using joins and subqueries
- Learn how to interact with a Postgres database from the command line

Demo and class exercise: Let's get up and running on Postgres!

Relational Database Management Systems (RDBMS)

A RDBMS is a type of database where data is stored in multiple related tables.

Example: A single table with records of customer purchases at an outdoor sports store.

Not very
efficient



| id | cust_name | cust_state | item_purchased | price | date |
|----|-----------|------------|----------------|-------|-------|
| 1 | John | CO | skis | \$300 | 10/30 |
| 2 | John | CO | goggles | \$75 | 11/14 |
| 3 | Taryn | CO | snowboard | \$400 | 11/18 |
| 4 | Adam | NY | skis | \$300 | 12/11 |
| 5 | Frank | AZ | skis | \$300 | 12/19 |
| 6 | Adam | NY | goggles | \$75 | 12/24 |

Relational Database Management Systems (RDBMS)

A RDBMS is a type of database where data is stored in multiple related tables.

Example: The same information in multiple tables in database.

customers

| cust_id | cust_name | cust_state |
|---------|-----------|------------|
| 1 | John | CO |
| 2 | Taryn | CO |
| 3 | Adam | NY |
| 4 | Frank | AZ |

products

| prod_id | description | price |
|---------|-------------|-------|
| 1 | skis | \$300 |
| 2 | goggles | \$75 |
| 3 | snowboard | \$400 |

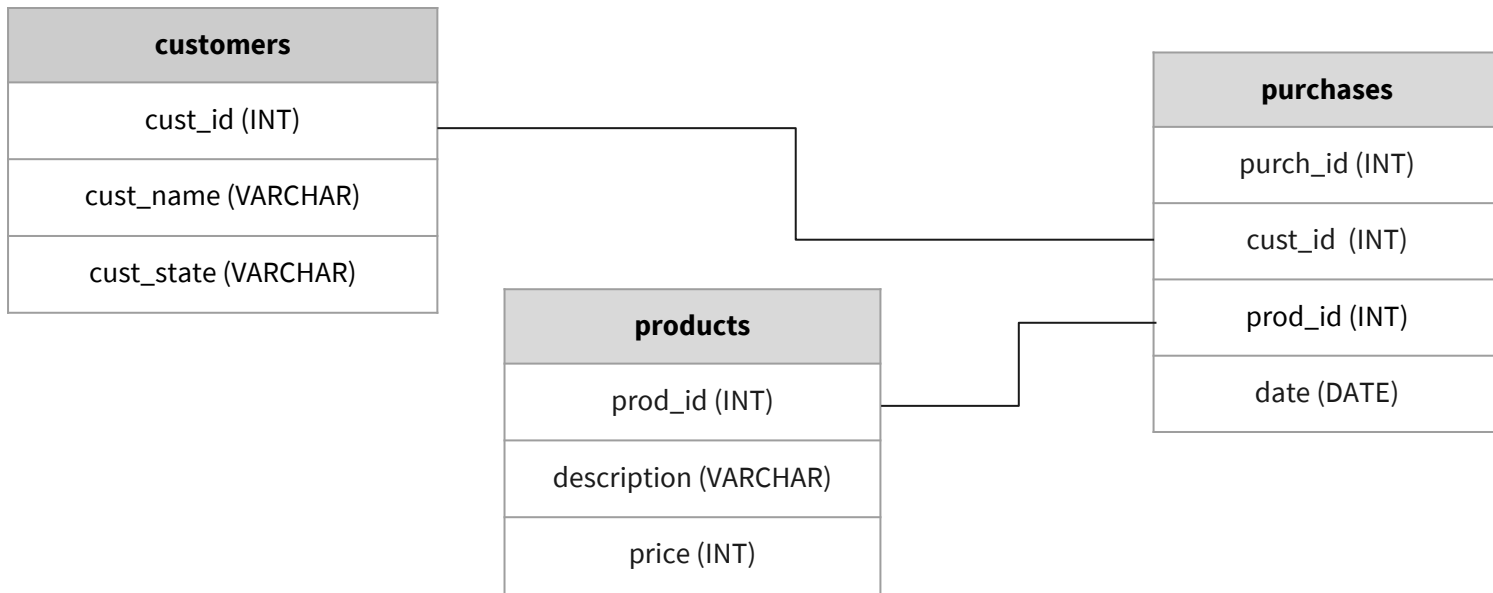
purchases

| cust_id | prod_id | date |
|---------|---------|-------|
| 1 | 1 | 10/30 |
| 1 | 2 | 11/14 |
| 2 | 3 | 11/18 |
| 3 | 1 | 12/11 |
| 4 | 1 | 12/19 |
| 3 | 2 | 12/24 |

Relational Database Management Systems (RDBMS)

A RDBMS is a type of database where data is stored in multiple related tables.

Example: The same information shown in an Entity Relationship Diagram (ERD):



Why RDBMS?



RDBMS provides one means of *persistent* data storage.

- Survives after the process in which it was created has ended
- Written to non-volatile storage (stored even if unpowered)
- Frequently accessed and unlikely to change in structure
- (e.g., a company database that contains records of customers and purchases)

Why RDBMS?



RDBMS provides to ability to:

- Model relations in data
- Query data and their relations efficiently
- Maintain data consistency and integrity

Why RDBMS?



For a long time, RDBMS was the *de facto* standard for storing data:

- Examples: Oracle, MySQL, SQL Server, Postgres
- In the era of “Big Data,” this is beginning to change
- But RDBMS are still everywhere and every data scientist should know how to work with them

RDBMS Terminology



- **Schema** defines the structure of a tables or a database
- Database is composed of a number of user-defined **tables**
- Each tables has **columns** (or fields) and **rows** (or records)
- A column is of a certain **data type** such as an integer, VARCHAR (str), or date

With a new data source, your first task is typically to understand the schema.

Always try to develop a holistic understanding of what you're looking at before diving into the details!

Structured Query Language (SQL)



SQL is the tool we use to interact with RDBMS. We can use SQL commands to:

- Create tables
- Alter tables
- Insert records
- Update records
- Delete records
- **Query (SELECT) records within or across tables**

The most critical skill for a Data Scientist--as opposed to a Data Engineer or Database Administrator--is to extract information from databases.

We will focus on writing queries in PostgreSQL, but all of the commands use similar vocabulary and syntax.

SQL Query Basics



All SQL queries have two main components:

```
SELECT    # What data (columns) do you want?
```

```
FROM      # From what location (table) you want it?
```

Note: SQL queries always return tables.

Note: SQL is a *declarative* language, unlike Python, which is *imperative*. With a declarative language, you tell the machine *what* you want, instead of *how*, and it figures out the best way to do it for you.

SELECT *

TABLE(S)

customers

| cust_id | cust_name | cust_state |
|---------|-----------|------------|
| 1 | John | CO |
| 2 | Taryn | CO |
| 3 | Adam | NY |
| 4 | Frank | AZ |

QUERY

```
SELECT  
    *  
FROM  
    customers;
```

OUTPUT

| cust_id | cust_name | cust_state |
|---------|-----------|------------|
| 1 | John | CO |
| 2 | Taryn | MA |
| 3 | Adam | NY |
| 4 | Frank | AZ |

The asterisk means “everything.”

Aliases

TABLE(S)

customers

| cust_id | cust_name | cust_state |
|---------|-----------|------------|
| 1 | John | CO |
| 2 | Taryn | CO |
| 3 | Adam | NY |
| 4 | Frank | AZ |

QUERY

```
SELECT
    cust_name AS name,
    cust_state state
FROM
    customers;
```

OUTPUT

| name | state |
|-------|-------|
| John | CO |
| Taryn | CO |
| Adam | NY |
| Frank | AZ |

- Aliasing can be used to rename columns and even tables (more on this later).
- “AS” makes code clearer but is not necessary.
- Be careful not to use keywords (e.g. count) as aliases!

Formatting SQL statements



Unlike Python, whitespace and capitalization do not matter (except for strings)

```
select column1, column2 from my_table;
```

Convention is to use ALL CAPS for keywords

Line breaks and indentation help make queries more readable (especially complex ones)

```
SELECT
    column1,
    column2
FROM
    My_table;
```

Punctuation such as commas (between items under each clause) and semicolons (after each statement) are required for proper evaluation

LIMIT and ORDER BY

TABLE(S)

customers

| cust_id | cust_name | cust_state |
|---------|-----------|------------|
| 1 | John | CO |
| 2 | Taryn | CO |
| 3 | Adam | NY |
| 4 | Frank | AZ |

QUERY

```
SELECT
    *
FROM
    customers
ORDER BY
    cust_name DESC
LIMIT 3;
```

OUTPUT

| cust_id | cust_name | cust_state |
|---------|-----------|------------|
| 1 | John | CO |
| 4 | Frank | AZ |
| 2 | Taryn | CO |

- ORDER BY is ascending by default; specify DESC for reverse sorting
- LIMIT specifies the number of records returned

SELECT DISTINCT

TABLE(S)

customers

| cust_id | cust_name | cust_state |
|---------|-----------|------------|
| 1 | John | CO |
| 2 | Taryn | CO |
| 3 | Adam | NY |
| 4 | Frank | AZ |

QUERY

```
SELECT DISTINCT
    cust_state
FROM
    customers;
```

OUTPUT

| cust_state |
|------------|
| CO |
| NY |
| AZ |

- SELECT DISTINCT grabs all the unique records.
- If multiple columns are selected, then all unique combinations are returned.

WHERE

TABLE(S)

customers

| cust_id | cust_name | cust_state |
|---------|-----------|------------|
| 1 | John | CO |
| 2 | Taryn | CO |
| 3 | Adam | NY |
| 4 | Frank | AZ |

QUERY

```
SELECT
    cust_name AS name,
    cust_state AS state
FROM
    customers
WHERE
    cust_state = 'CO' ;
```

OUTPUT

| name | state |
|-------|-------|
| John | CO |
| Taryn | CO |

- WHERE specifies criterion for selecting specific rows (row filter)
- Note that the WHERE statement must reference the original column name, not the alias
- However, WHERE can reference a table column that is not in SELECT (e.g. cust_id)

WHERE (Multiple Criteria)

TABLE(S)

customers

| cust_id | cust_name | cust_state |
|---------|-----------|------------|
| 1 | John | CO |
| 2 | Taryn | CO |
| 3 | Adam | NY |
| 4 | Frank | AZ |

QUERY

```
SELECT
    cust_name AS name,
    cust_state AS state
FROM
    customers
WHERE
    (cust_state = 'CO'
    AND cust_name = 'John')
    OR cust_state = 'NY' ;
```

OUTPUT

| name | state |
|------|-------|
| John | CO |
| Adam | NY |

- We can specify multiple conditions on the “WHERE” clause by using AND/OR
- Note that comparison operator uses a single equal sign (= instead of ==)

ARITHMETIC OPERATORS (+, -, *, /, etc.)

TABLE(S)

products

| prod_id | description | price |
|---------|-------------|-------|
| 1 | skis | 300 |
| 2 | goggles | 75 |
| 3 | snowboard | 400 |

QUERY

```
SELECT
    description,
    price,
    price * 2 AS ripoff
FROM
    products;
```

OUTPUT

| description | price | ripoff |
|-------------|-------|--------|
| skis | 300 | 600 |
| goggles | 75 | 150 |
| snowboard | 400 | 800 |

- Arithmetic operators are similar to Python (except SQL uses ^ for exponents)
- Can be used with multiple columns (for example, adding one column value to another)

ARITHMETIC OPERATORS and DATA TYPES

TABLE(S)

QUERY

OUTPUT

products

| prod_id | description | price |
|---------|-------------|-------|
| 1 | skis | 300 |
| 2 | goggles | 75 |
| 3 | snowboard | 400 |

```
SELECT
    description,
    price,
    price/2 AS sale_int,
    price/2. AS sale_float
FROM
    products;
```

| description | price | sale_int | sale_float |
|-------------|-------|----------|------------|
| skis | 300 | 150 | 150.0 |
| goggles | 75 | 37 | 37.5 |
| snowboard | 400 | 200 | 200.0 |

- Arithmetic operators are similar to Python (except SQL uses ^ for exponents)
- Can be used with multiple columns (for example, adding one column value to another)

BREAKOUT!



pets

| id | name | species | age | gender | owner |
|----|--------|---------|-----|--------|-------|
| 1 | Max | cat | 8 | M | Taryn |
| 2 | Belle | cat | 10 | F | Taryn |
| 3 | Bailey | dog | 11 | F | Kyrie |
| 4 | Daisy | cat | 5 | F | Kyrie |
| 5 | Kahlua | dog | 7 | F | Blair |
| 6 | Henley | dog | 9 | F | Megan |
| 7 | Salem | cat | 1 | F | Megan |
| 8 | Teeny | cat | 1 | F | Megan |

Write queries that would return:

- 1) Owner(s) of Male pet(s)
- 2) Names of dogs
- 3) Names and ages of oldest two pets
- 4) The species of the youngest pet
- 5) Names of cats that are 8 years old or younger
- 6) Pets that are babies (≤ 1 year) are expensive. Senior pets (≥ 8) can also be expensive. Who owns one or more expensive pets?

Pick-me-up before we return to regular programming....



CASE WHEN

TABLE(S)

customers

| cust_id | cust_name | cust_state |
|---------|-----------|------------|
| 1 | John | CO |
| 2 | Taryn | CO |
| 3 | Adam | NY |
| 4 | Frank | AZ |

QUERY

```
SELECT
    cust_name AS name,
    CASE WHEN cust_state = 'CO' THEN 1
          ELSE 0 END AS in_state
FROM
    customers;
```

OUTPUT

| name | in_state |
|-------|----------|
| John | 1 |
| Taryn | 1 |
| Adam | 0 |
| Frank | 0 |

- CASE WHEN statement is the SQL version of an if-then-else statement
- Used in the SELECT clause
- Can combine multiple WHEN statements and/or multiple conditionals

Aggregators

TABLE(S)

products

| prod_id | description | price |
|---------|-------------|-------|
| 1 | skis | 300 |
| 2 | goggles | 75 |
| 3 | snowboard | 400 |

QUERY

```
SELECT
    COUNT(*) ,
    MAX(price)
FROM
    products;
```

OUTPUT

| COUNT | MAX |
|-------|-----|
| 3 | 400 |

- Aggregators combine information from multiple rows into a single row.
- Other aggregators include MIN, MAX, SUM, COUNT, STDDEV, etc.

GROUP BY

TABLE(S)

customers

| cust_id | cust_name | cust_state |
|---------|-----------|------------|
| 1 | John | CO |
| 2 | Taryn | CO |
| 3 | Adam | NY |
| 4 | Frank | AZ |

QUERY

```
SELECT
    cust_state as state,
    count(*)
FROM
    customers
GROUP BY
    cust_state;
```

OUTPUT

| state | count(*) |
|-------|----------|
| CO | 2 |
| NY | 1 |
| AZ | 1 |

- The GROUP BY clause calculates aggregate statistics for groups of data
- **Any column that is not an aggregator *must be in the GROUP BY clause*** (for example, if we added `cust_name` to the SELECT clause only, SQL would not know whether to return John or Taryn in the CO row)
- Any column in the GROUP BY by clause must also appear in the SELECT clause (true of Postgres but not MySQL or Spark SQL)

GROUP BY and WHERE

TABLE(S)

customers

| cust_id | cust_name | cust_state |
|---------|-----------|------------|
| 1 | John | CO |
| 2 | Taryn | CO |
| 3 | Adam | NY |
| 4 | Frank | AZ |

QUERY

```
SELECT
    cust_state AS state,
    COUNT(*) AS total
FROM
    customers
WHERE
    cust_name != 'Adam'
GROUP BY
    cust_state;
```

OUTPUT

| state | total |
|-------|-------|
| CO | 2 |
| AZ | 1 |

GROUP BY and WHERE (cont'd)

TABLE(S)

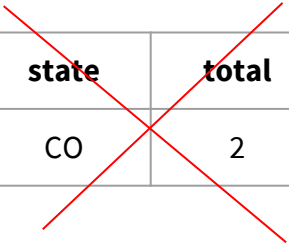
customers

| cust_id | cust_name | cust_state |
|---------|-----------|------------|
| 1 | John | CO |
| 2 | Taryn | CO |
| 3 | Adam | NY |
| 4 | Frank | AZ |

QUERY

```
SELECT
    cust_state AS state,
    COUNT(*) AS total
FROM
    customers
WHERE
    COUNT(*) >= 2
GROUP BY
    cust_state;
```

OUTPUT



| state | total |
|-------|-------|
| CO | 2 |

ERROR

- Why does the query above not work?

GROUP BY and HAVING

TABLE(S)

customers

| cust_id | cust_name | cust_state |
|---------|-----------|------------|
| 1 | John | CO |
| 2 | Taryn | CO |
| 3 | Adam | NY |
| 4 | Frank | AZ |

QUERY

```
SELECT
    cust_state AS state,
    COUNT(*) AS total
FROM
    customers
WHERE
    count(*) >= 2
GROUP BY
    cust_state
HAVING
    COUNT(*) >= 2;
```

OUTPUT

| state | total |
|-------|-------|
| CO | 2 |

- Use HAVING instead of WHERE when filtering rows *after* aggregation
- WHERE clause filters rows in the root table *before* aggregation
- Like WHERE clause, HAVING clause cannot reference an alias (in Postgres, at least)

Joining Tables



The JOIN clause allows us to use a single query to extract information from multiple tables.

Every JOIN statement has two parts:

1. Specifying the tables to be joined (JOIN)
2. Specifying the columns to join tables on (ON)

For example, we could learn the home state of every purchaser of an item:

1. JOIN the *purchases* table (history of purchase events) and the *customers* table (info about customers)
2. ON the *cust_id* column, which appears in both tables

Primary Keys



- Every table in a RDBMS has a **primary key (PK)** that uniquely identifies that row
- Each entry must have a PK, and PKs cannot repeat within a table
- PKs are usually integers but can take other forms

customers

| cust_id | cust_name | cust_state |
|---------|-----------|------------|
| 1 | John | CO |
| 2 | Taryn | CO |
| 3 | Adam | NY |
| 4 | Frank | AZ |

products

| prod_id | description | price |
|---------|-------------|-------|
| 1 | skis | \$300 |
| 2 | goggles | \$75 |
| 3 | snowboard | \$400 |

purchases

| purch_id | cust_id | prod_id | date |
|----------|---------|---------|-------|
| 1 | 1 | 1 | 10/30 |
| 2 | 1 | 2 | 11/14 |
| 3 | 2 | 3 | 11/18 |
| 4 | 3 | 1 | 12/11 |
| 5 | 4 | 1 | 12/19 |
| 6 | 3 | 2 | 12/24 |

Foreign Keys and Table Relationships



- A foreign key (FK) is a column that uniquely identifies a column in another table
- Often, a FK in one table is a PK in another table (but not necessarily)
- We can use FKs to join tables

customers

| cust_id | cust_name | cust_state |
|---------|-----------|------------|
| 1 | John | CO |
| 2 | Taryn | CO |
| 3 | Adam | NY |
| 4 | Frank | AZ |

products

| prod_id | description | price |
|---------|-------------|-------|
| 1 | skis | \$300 |
| 2 | goggles | \$75 |
| 3 | snowboard | \$400 |

purchases

| purch_id | cust_id | prod_id | date |
|----------|---------|---------|-------|
| 1 | 1 | 1 | 10/30 |
| 2 | 1 | 2 | 11/14 |
| 3 | 2 | 3 | 11/18 |
| 4 | 3 | 1 | 12/11 |
| 5 | 4 | 1 | 12/19 |
| 6 | 3 | 2 | 12/24 |

Relationship Types



Foreign keys models a few different types of relationships:

- **One-to-many:** cust_id and purch_id
- **Many-to-many:** cust_id and prod_id
- **One-to-one:** sku_id and prod_id

customers

| cust_id | cust_name | cust_state |
|---------|-----------|------------|
| 1 | John | CO |
| 2 | Taryn | CO |
| 3 | Adam | NY |
| 4 | Frank | AZ |

purchases

| purch_id | cust_id | prod_id | date |
|----------|---------|---------|-------|
| 1 | 1 | 1 | 10/30 |
| 2 | 1 | 2 | 11/14 |
| 3 | 2 | 3 | 11/18 |
| 4 | 3 | 1 | 12/11 |
| 5 | 4 | 1 | 12/19 |
| 6 | 3 | 2 | 12/24 |

products

| prod_id | description | price |
|---------|-------------|-------|
| 1 | skis | \$300 |
| 2 | goggles | \$75 |
| 3 | snowboard | \$400 |

product_SKUs

| sku_id | prod_id |
|----------|---------|
| 1413434 | 1 |
| 7587578 | 2 |
| 35635635 | 3 |

purchases

| purch_id | cust_id | prod_id | date |
|----------|---------|---------|-------|
| 1 | 1 | 1 | 10/30 |
| 2 | 1 | 2 | 11/14 |
| 3 | 2 | 3 | 11/18 |
| 4 | 3 | 1 | 12/11 |
| 5 | 4 | 1 | 12/19 |
| 6 | 3 | 2 | 12/24 |

customers

| cust_id | cust_name | cust_state |
|---------|-----------|------------|
| 1 | John | CO |
| 2 | Taryn | CO |
| 3 | Adam | NY |
| 4 | Frank | AZ |

JOINS

QUERY

```
SELECT
    purchases.purch_id,
    customers.cust_id,
    customers.cust_state
FROM
    purchases
JOIN
    customers
ON
    purchases.cust_id =
    customers.cust_id;
```




OUTPUT

| purch_id | cust_id | cust_state |
|----------|---------|------------|
| 1 | 1 | CO |
| 2 | 1 | CO |
| 3 | 2 | CO |
| 4 | 3 | NY |
| 5 | 4 | AZ |
| 6 | 3 | NY |

JOIN Types

SELECT ... FROM TableA _____ JOIN
TableB ON...

- **(INNER) JOIN:** Discards any entries that do not have match between the keys specified in the ON clause
- **LEFT (OUTER) JOIN:** Keeps all entries in the left (FROM) table, regardless of whether any matches are found in the right (JOIN) tables
 - **RIGHT (OUTER) JOIN:** Is the same, except keeps all entries in the right (JOIN) table instead of the left (FROM) table); usually avoided because it does the same thing as a LEFT join
- **FULL (OUTER) JOIN:** Keeps the rows in both tables no matter what



purchases

| purch_id | cust_id | prod_id | date |
|----------|---------|---------|-------|
| 1 | 1 | 1 | 10/30 |
| 2 | 1 | 2 | 11/14 |
| 3 | 2 | 3 | 11/18 |
| 4 | 3 | 1 | 12/11 |
| 5 | NULL | 1 | 12/19 |
| 6 | NULL | 2 | 12/24 |

customers

| cust_id | cust_name | cust_state |
|---------|-----------|------------|
| 1 | John | CO |
| 2 | Taryn | CO |
| 3 | Adam | NY |
| 4 | Frank | AZ |
| 5 | Neil | NY |

(INNER) JOIN

QUERY

```
SELECT
    purchases.purch_id,
    customers.cust_id,
    customers.cust_state
FROM
    purchases
INNER JOIN
    customers
ON
    purchases.cust_id =
    customers.cust_id;
```

OUTPUT

| purch_id | cust_id | cust_state |
|----------|---------|------------|
| 1 | 1 | CO |
| 2 | 1 | CO |
| 3 | 2 | CO |
| 4 | 3 | NY |

INNER JOIN discards records that do not have a match in both tables

purchases

| purch_id | cust_id | prod_id | date |
|----------|---------|---------|-------|
| 1 | 1 | 1 | 10/30 |
| 2 | 1 | 2 | 11/14 |
| 3 | 2 | 3 | 11/18 |
| 4 | 3 | 1 | 12/11 |
| 5 | NULL | 1 | 12/19 |
| 6 | NULL | 2 | 12/24 |

customers

| cust_id | cust_name | cust_state |
|---------|-----------|------------|
| 1 | John | CO |
| 2 | Taryn | CO |
| 3 | Adam | NY |
| 4 | Frank | AZ |
| 5 | Neil | NY |

LEFT (OUTER) JOIN

QUERY

```
SELECT
    purchases.purch_id,
    customers.cust_id,
    customers.cust_state
FROM
    purchases
LEFT OUTER JOIN
    customers
ON
    purchases.cust_id =
    customers.cust_id;
```

OUTPUT

| purch_id | cust_id | cust_state |
|----------|---------|------------|
| 1 | 1 | CO |
| 2 | 1 | CO |
| 3 | 2 | CO |
| 4 | 3 | NY |
| 5 | NULL | NULL |
| 6 | NULL | NULL |

LEFT OUTER JOIN retains all records from the left (FROM) tables and includes records from the right (JOIN) table if they are available

purchases

| purch_id | cust_id | prod_id | date |
|----------|---------|---------|-------|
| 1 | 1 | 1 | 10/30 |
| 2 | 1 | 2 | 11/14 |
| 3 | 2 | 3 | 11/18 |
| 4 | 3 | 1 | 12/11 |
| 5 | NULL | 1 | 12/19 |
| 6 | NULL | 2 | 12/24 |

customers

| cust_id | cust_name | cust_state |
|---------|-----------|------------|
| 1 | John | CO |
| 2 | Taryn | CO |
| 3 | Adam | NY |
| 4 | Frank | AZ |
| 5 | Neil | NY |

FULL (OUTER) JOIN

QUERY

```
SELECT
    purchases.purch_id,
    customers.cust_id,
    customers.cust_state
FROM
    purchases
FULL OUTER JOIN
    customers
ON
    purchases.cust_id =
    customers.cust_id;
```

OUTPUT

| purch_id | cust_id | cust_state |
|----------|---------|------------|
| 1 | 1 | CO |
| 2 | 1 | CO |
| 3 | 2 | CO |
| 4 | 3 | NY |
| 5 | NULL | NULL |
| 6 | NULL | NULL |
| NULL | 4 | AZ |
| NULL | 5 | NY |

FULL OUTER JOIN retains all records from both tables regardless of matches

Query Components vs. Order of Evaluation



1. FROM + JOIN: first the product of all tables is formed
2. WHERE: the where clause filters rows that do not meet the search condition
3. GROUP BY + (COUNT, SUM, etc): the rows are grouped using the columns in the group by clause and the aggregation functions are applied on the grouping
4. HAVING: like the WHERE clause, but can be applied after aggregation
5. SELECT: the targeted list of columns are evaluated and returned
6. DISTINCT: duplicate rows are eliminated
7. ORDER BY: the resulting rows are sorted

Order of Evaluation - subtle points



WHERE clause: eliminate rows you don't want early. EFFICIENCY! Even more important in when querying from distributed file systems

WHERE and GROUP BY are evaluated *before* SELECT statement. If you do an aggregation/ name change/other manipulation, you will need to use the original column name here because your new alias won't be recognized.

ORDER BY is evaluated *after* the SELECT statement. Use new aliases

Subqueries



- In general, you can replace any table name with a subquery:

```
SELECT ... FROM (SELECT ...)
```

- If a query returns a single value, you can use it as such:

```
...WHERE column1 = (SELECT ...)
```


- If a query returns a single column, you can treat it like a vector:

```
...WHERE column1 IN (SELECT ...)
```

Using Postgres from the Command Line

- Instructions on Postgres installation and set-up are in the *individual.md* file
- Postgres must be running in order to use it from the command line:

Look for
ambiguous
ungulate icon



- Instructions on loading the database and entering postgres prompt from the command line are also in the *individual.md* file

Load .sql file into a DB and run queries



One-time step to create a database and load .sql file. From the command line:

```
psql
CREATE DATABASE MyDatabase;    -- whatever name you choose
\q
psql MyDatabase < file.sql
```

Now you can access this database any time:

```
psql MyDataBase
```

Using Postgres from the Command Line (cont'd)



Useful commands from the psql interactive shell prompt:

- `\l` - list all databases
- `\d` - list all tables
- `\d <table name>` - describe a table's schema
- `\h <clause>` - Help for SQL clause help
- `q` - exit current view and return to command line
- `\q` - quit psql
- `\i script.sql` - run script (or query)

Morning Objectives



- Discuss RDBMS and why we use them
- Write simple SQL queries on single table using SELECT, FROM, WHERE, GROUP BY, ORDER BY clauses as well as aggregation functions (COUNT, AVG, etc.)
- Understand primary keys, foreign keys, and table relationships
- Write complex queries using joins and subqueries
- Learn how to interact with a Postgres database from the command line

Demo and class exercise: Let's get up and running on Postgres!